

D Bippin Shekar Goud

9100863981 | bippinshekar@gmail.com | [LinkedIn](#) | [GitHub](#)

EXPERIENCE

Founder & Lead Engineer

Dec 2025 – Present

Leaps

Early-stage

- Building an AI-native career execution platform that scores candidate fit, identifies skill gaps, and autonomously applies to high-match roles, replacing volume-based job search with fit-driven guided execution.
- Engineered autonomous application systems for LinkedIn (multi-step, turn-based, dynamically structured per company) and Ashby (single-form with company-specific open-ended fields requiring contextual candidate reasoning), two fundamentally different automation problems resolved under one architecture.
- Designed the full multi-agent system covering resume parsing, fit scoring, skill gap analysis, and stateful execution agents that navigate dynamic application flows end to end.

Founding AI Engineer

Nov 2025 – Jan 2026

Pulsegen.io

- Rebuilt a stateless, prompt-based analytics bot into a full agentic TypeScript system over unstructured MongoDB, introducing Text-to-Mongo query generation and a contextual memory layer the original had none of, cutting end-to-end latency 40% (80s to 40–50s) and improving retrieval accuracy and answer consistency.
- Scaled an LLM-driven VoC engine that turned unstructured customer feedback into decision-ready product insights, increasing throughput 33x and handling 100K+ inputs/day in production.

Founding AI Engineer

Sept 2024 – Nov 2025

8bit.ai

- Built an enterprise Text-to-SQL engine from scratch across PostgreSQL, MySQL, and a data lake (Iceberg) spanning 200+ tables across multiple schemas, including a metadata jargon mapping layer that translated business terminology (e.g. "cumulative ARR by department") into actual column structures without users needing to know data team naming conventions.
- Fine-tuned Llama 70B (QLoRA) and OpenAI's fine-tunable model on 60,000+ curated business SQL pairs, raising SQL generation from 80% to 98% and execution success from 55% to 85%+, with pre/post guardrail and SQL correction layers accounting for the remaining production reliability.
- Designed conversational orchestration routing across T2SQL, RAG, workflow generation, and web research, with Redis-backed hierarchical memory classification replacing rolling context windows to maintain coherent reasoning across long conversations.
- Migrated orchestration from LangChain to LangGraph for HITL flows, then eliminated all framework dependencies in favor of direct SDK integrations with custom provider abstractions, removing black-box behavior entirely.
- Led Go backend development for the AI analytics dashboard with a teammate and managed all self-hosted Kubernetes deployments (Rancher, then direct K8s with Headlamp) with Docker throughout.

PROJECTS

Askrr | *Open Source, In Development* | [GitHub](#)

- Building an open-source tool for engineers to optimize their AI-assisted development cycles, built out of personal need for better day-to-day development tooling.

TECHNICAL SKILLS

Languages: Python, Go, TypeScript, SQL, Bash

Backend & Infrastructure: FastAPI, Go Services, Redis, Docker, Kubernetes, REST APIs, GraphQL

Frontend: Next.js, React

AI / ML: LLMs, OpenAI, Gemini, Llama, PyTorch, TensorFlow

LLM Systems: RAG, Fine-tuning (QLoRA, LoRA), Prompt Engineering, Evaluation, Guardrails, HITL Architectures

Vector & Retrieval: FAISS, Pinecone, Weaviate, Hybrid Search, Semantic Search, Embeddings

Agents & Orchestration: Multi-Agent Systems, LangGraph, Tool Calling, Conversational Orchestration

Databases: PostgreSQL, MySQL, MongoDB, Redis, Data Lakes (Iceberg)

EDUCATION

St Mary's Group of Institutions

Hyderabad, India

B.Tech. in Computer Science, Artificial Intelligence

2020 – 2025